

In [1]:

```
import pandas as pd
import numpy as np

#####
#
#   Decision Tree, Random Forest & Gradient Boosting
#
#####

#####
#>>> 1> Read file

titanic = pd.read_csv('http://biostat.mc.vanderbilt.edu/wiki/pub/Main/DataSets/titanic.txt')

print(titanic.head())#Pre-look at titanic as a dataframe
```

	row.names	pclass	survived	\
0	1	1st	1	
1	2	1st	0	
2	3	1st	0	
3	4	1st	0	
4	5	1st	1	

		name	age	embarked	\
0		Allen, Miss Elisabeth Walton	29.0000	Southampton	
1		Allison, Miss Helen Loraine	2.0000	Southampton	
2		Allison, Mr Hudson Joshua Creighton	30.0000	Southampton	
3	Allison, Mrs Hudson J.C. (Bessie Waldo Daniels)		25.0000	Southampton	
4		Allison, Master Hudson Trevor	0.9167	Southampton	

		home.dest	room	ticket	boat	sex
0		St Louis, MO	B-5	24160 L221	2	female
1	Montreal, PQ / Chesterville, ON		C26	NaN	NaN	female
2	Montreal, PQ / Chesterville, ON		C26	NaN	(135)	male
3	Montreal, PQ / Chesterville, ON		C26	NaN	NaN	female
4	Montreal, PQ / Chesterville, ON		C22	NaN	11	male

In [2]:

```
print(titanic.info())#Look at titanic statistic information
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1313 entries, 0 to 1312
Data columns (total 11 columns):
row.names      1313 non-null int64
pclass         1313 non-null object
survived       1313 non-null int64
name           1313 non-null object
age            633 non-null float64
embarked       821 non-null object
home.dest      754 non-null object
room           77 non-null object
ticket         69 non-null object
boat           347 non-null object
sex            1313 non-null object
dtypes: float64(1), int64(2), object(8)
memory usage: 112.9+ KB
None
```

In [3]:

```
#####
#>>> 2> Select features and Vectorize
```

```
X = titanic[['pclass', 'age', 'sex']]
y = titanic['survived']
```

```
print(X.info()) # Check features information
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1313 entries, 0 to 1312
Data columns (total 3 columns):
pclass      1313 non-null object
age         633 non-null float64
sex         1313 non-null object
dtypes: float64(1), object(2)
memory usage: 30.9+ KB
None
```

In [4]:

```
X['age'].fillna(X['age'].mean(), inplace=True)
print(X.info()) # Check again after filling out
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1313 entries, 0 to 1312
Data columns (total 3 columns):
pclass    1313 non-null object
age        1313 non-null float64
sex        1313 non-null object
dtypes: float64(1), object(2)
memory usage: 30.9+ KB
None
```

C:\Program Files\Anaconda3\lib\site-packages\pandas\core\generic.py:3191: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

```
self._update_inplace(new_data)
```

In [5]:

```
from sklearn.cross_validation import train_test_split
from sklearn.feature_extraction import DictVectorizer
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=33) #Split data
```

```
vec = DictVectorizer(sparse=False) #Vectorize
X_train=vec.fit_transform(X_train.to_dict(orient='record'))
print(vec.feature_names_)
X_test=vec.transform(X_test.to_dict(orient='record'))
```

```
['age', 'pclass=1st', 'pclass=2nd', 'pclass=3rd', 'sex=female', 'sex=male']
```

In [7]:

```
#####  
#>>> 3> Build model and Predict  
  
# Decision Tree  
from sklearn.tree import DecisionTreeClassifier  
  
dtc=DecisionTreeClassifier()  
dtc.fit(X_train,y_train)  
y_predict=dtc.predict(X_test)  
  
# Random Forest  
from sklearn.ensemble import RandomForestClassifier  
rfc=RandomForestClassifier()  
rfc.fit(X_train,y_train)  
rfc_y_pred=rfc.predict(X_test)  
  
# Gradient Boosting  
from sklearn.ensemble import GradientBoostingClassifier  
gbc=GradientBoostingClassifier()  
gbc.fit(X_train,y_train)  
gbc_y_pred=gbc.predict(X_test)  
  
#####  
#>>> 4> Score  
from sklearn.metrics import classification_report  
  
# Decision Tree  
print('The accuracy of Decision Tree is',dtc.score(X_test,y_test))  
print(classification_report(y_predict,y_test))  
  
# Random Forest  
print('The accuracy of Random Forest Classifier is',rfc.score(X_test,y_test))  
print(classification_report(rfc_y_pred,y_test))  
  
# Gradient Boosting  
print('The accuracy of Gradient Boosting Classifier is',gbc.score(X_test,y_test))  
print(classification_report(gbc_y_pred,y_test))
```

The accuracy of Decision Tree is 0.781155015198

	precision	recall	f1-score	support
0	0.91	0.78	0.84	236
1	0.58	0.80	0.67	93
avg / total	0.81	0.78	0.79	329

The accuracy of Random Forest Classifier is 0.781155015198

	precision	recall	f1-score	support
0	0.90	0.78	0.83	234
1	0.59	0.79	0.68	95
avg / total	0.81	0.78	0.79	329

The accuracy of Gradient Boosting Classifier is 0.790273556231

	precision	recall	f1-score	support
0	0.92	0.78	0.84	239
1	0.58	0.82	0.68	90
avg / total	0.83	0.79	0.80	329

In [8]:

```
#####  
#  
#   Logistic Regression & SGD Regression  
#  
#####  
  
#####  
#>>> 1> Read file  
  
from sklearn.datasets import load_boston  
boston=load_boston()  
print(boston.DESCR)
```

Boston House Prices dataset

Notes

Data Set Characteristics:

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive

:Median Value (attribute 14) is usually the target

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<http://archive.ics.uci.edu/ml/datasets/Housing> (<http://archive.ics.uci.edu/ml/datasets/Housing>)

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.
- many more! (see <http://archive.ics.uci.edu/ml/datasets/Housing>) (<http://archive.ics.uci.edu/ml/datasets/Housing>)

In [9]:

```
#####  
#>>> 2> Split data and preprocessing  
  
# Split data  
from sklearn.cross_validation import train_test_split  
  
X=boston.data  
y=boston.target  
  
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.25,random_state=33) #Split data  
print('The max target value is', np.max(boston.target))  
print('The min target value is', np.min(boston.target))  
print('The average target value is', np.mean(boston.target))  
  
# preprocessing  
from sklearn.preprocessing import StandardScaler  
  
ss_X=StandardScaler()  
ss_y=StandardScaler()  
  
X_train=ss_X.fit_transform(X_train)  
X_test=ss_X.fit_transform(X_test)  
  
y_train=ss_y.fit_transform(y_train)  
y_test=ss_y.fit_transform(y_test)
```

The max target value is 50.0

The min target value is 5.0

The average target value is 22.5328063241

C:\Program Files\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:583: DeprecationWarning: Passing 1d arrays as data is deprecated in 0.17 and will raise ValueError in 0.19. Reshape your data either using X.reshape(-1, 1) if your data has a single feature or X.reshape(1, -1) if it contains a single sample.

warnings.warn(DEPRECATION_MSG_1D, DeprecationWarning)

C:\Program Files\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:646: DeprecationWarning: Passing 1d arrays as data is deprecated in 0.17 and will raise ValueError in 0.19. Reshape your data either using X.reshape(-1, 1) if your data has a single feature or X.reshape(1, -1) if it contains a single sample.

warnings.warn(DEPRECATION_MSG_1D, DeprecationWarning)

C:\Program Files\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:583: DeprecationWarning: Passing 1d arrays as data is deprecated in 0.17 and will raise ValueError in 0.19. Reshape your data either using X.reshape(-1, 1) if your data has a single feature or X.reshape(1, -1) if it contains a single sample.

warnings.warn(DEPRECATION_MSG_1D, DeprecationWarning)

C:\Program Files\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:646: DeprecationWarning: Passing 1d arrays as data is deprecated in 0.17 and will raise ValueError in 0.19. Reshape your data either using X.reshape(-1, 1) if your data has a single feature or X.reshape(1, -1) if it contains a single sample.

warnings.warn(DEPRECATION_MSG_1D, DeprecationWarning)

In [10]:

```
#####  
#>>> 3> Build model and Predict  
  
# Linear Regression  
from sklearn.linear_model import LinearRegression  
  
lr=LinearRegression()  
lr.fit(X_train,y_train)  
lr_y_predict=lr.predict(X_test)  
  
# SGDRegressor  
from sklearn.linear_model import SGDRegressor  
sgdr=SGDRegressor()  
sgdr.fit(X_train,y_train)  
sgdr_y_pred=sgdr.predict(X_test)  
  
#####  
#>>> 4> Score  
  
# Linear Regression  
print('The value of default measurement of LinearRegression is',lr.score(X_test,y_test))  
  
from sklearn.metrics import r2_score,mean_squared_error,mean_absolute_error  
  
# R-squared  
print('The value of R-squared of LinearRegression is',r2_score(y_test,lr_y_predict))  
  
# Mean squared error  
print('The value of mean squared error of LinearRegression is',mean_squared_error(ss_y.inverse_trans  
  
# Mean absolute error  
print('The value of mean absolute error of LinearRegression is',mean_absolute_error(ss_y.inverse_tra
```

The value of default measurement of LinearRegression is 0.676930350524
The value of R-squared of LinearRegression is 0.676930350524
The value of mean squared error of LinearRegression is 25.0512388542
The value of mean absolute error of LinearRegression is 3.51371562676

In [45]:

```
#####  
#  
#          SVM  
#  
#####  
  
#####  
#>>> 1> Read file  
  
titanic = pd.read_csv('http://biostat.mc.vanderbilt.edu/wiki/pub/Main/DataSets/titanic.txt')  
  
#####  
#>>> 2> Vectorize  
  
X = titanic.drop(['name', 'row.names', 'survived'], axis=1)  
y = titanic['survived']  
  
#fill missing value  
X['age'].fillna(X['age'].mean(), inplace=True)  
X.fillna('UNKNOWN', inplace=True)  
  
from sklearn.cross_validation import train_test_split  
from sklearn.feature_extraction import DictVectorizer  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=33) #Split data  
  
vec = DictVectorizer()  
  
X_train=vec.fit_transform(X_train.to_dict(orient='record'))  
print(len(vec.feature_names_))  
X_test=vec.transform(X_test.to_dict(orient='record'))
```

474

In [46]:

```
#####  
#>>> 3> Build model  
  
# Decision Tree  
from sklearn.tree import DecisionTreeClassifier  
  
dt=DecisionTreeClassifier(criterion='entropy')  
dt.fit(X_train, y_train)  
print(dt.score(X_test, y_test))
```

0.823708206687

In [13]:

```
#####
#>>> 4> Select feature and predict
from sklearn import feature_selection

fs=feature_selection.SelectPercentile(feature_selection.chi2,percentile=20)

X_train_fs=fs.fit_transform(X_train,y_train)
dt.fit(X_train_fs,y_train)

X_test_fs=fs.transform(X_test)
print(dt.score(X_test_fs,y_test))

0.817629179331
```

In [47]:

```
#####
#>>> 5> Validation

from sklearn.cross_validation import cross_val_score
percentiles=range(1,100,2)
results=[]
for i in percentiles:
    fs=feature_selection.SelectPercentile(feature_selection.chi2,percentile=i)
    X_train_fs=fs.fit_transform(X_train,y_train)
    scores=cross_val_score(dt,X_train_fs,y_train,cv=5)
    results=np.append(results,scores.mean())
print(results)

[ 0.85063904  0.85673057  0.87501546  0.88622964  0.86590394  0.86998557
 0.86896516  0.87302618  0.86591424  0.87097506  0.86589363  0.86691404
 0.86795506  0.86386312  0.87097506  0.86590394  0.87199546  0.86691404
 0.86486291  0.87198516  0.8597918  0.86791383  0.86792414  0.87096475
 0.86894455  0.87199546  0.87097506  0.87404659  0.87198516  0.86893424
 0.86996496  0.86385281  0.87098536  0.86998557  0.87096475  0.86996496
 0.86688312  0.87201608  0.86995465  0.86795506  0.87098536  0.86691404
 0.86589363  0.86082251  0.86386312  0.8628221  0.86084313  0.8618223
 0.8598021  0.85675119]
```

In [48]:


```
opt=np.where(results==results.max())[0]
print('Optimal number of features %d'%percentiles[opt])#I still couldn't find out the reason why it
```

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-48-3fa7ef434f1d> in <module>()
      1 opt=np.where(results==results.max())[0]
----> 2 print('Optimal number of features %d'%percentiles[opt])
```

TypeError: only integer scalar arrays can be converted to a scalar index

In [31]:

```
#####  
#>>> 6> Visualization  
  
import pylab as pl  
pl.plot(percentiles, results)  
pl.xlabel('percentiles of features')  
pl.ylabel('accuracy')  
pl.show()
```



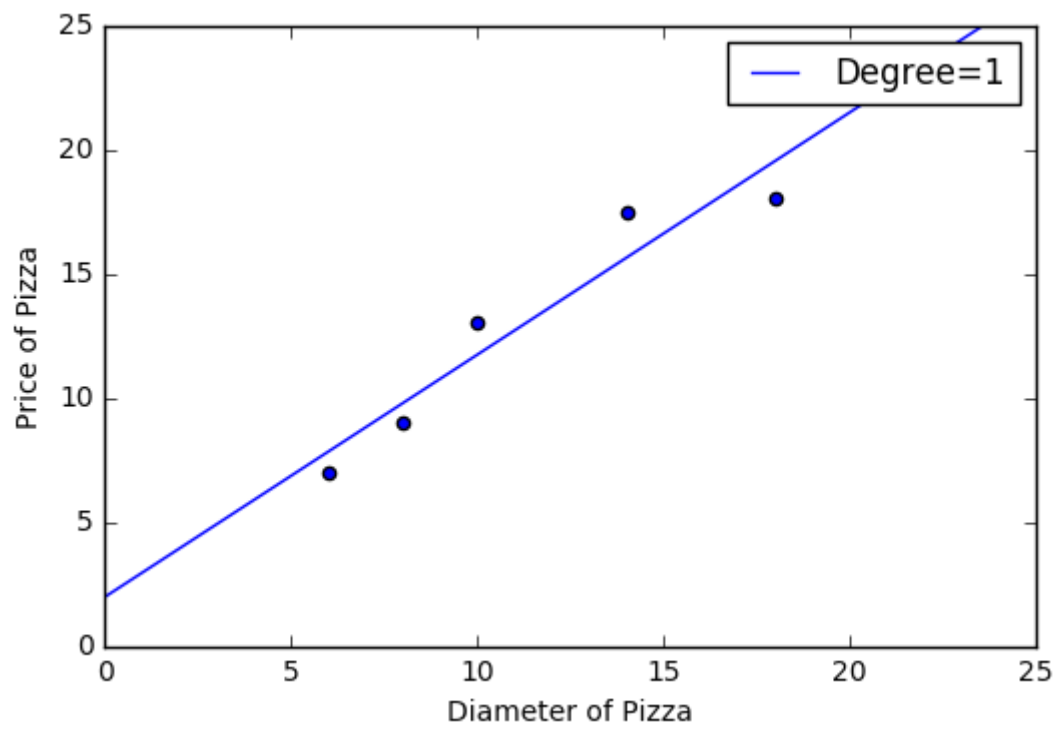
In [34]:

```
#Validation  
from sklearn import feature_selection  
  
fs=feature_selection.SelectPercentile(feature_selection.chi2, percentile=20)  
  
X_train_fs=fs.fit_transform(X_train, y_train)  
dt.fit(X_train_fs, y_train)  
  
X_test_fs=fs.transform(X_test)  
print(dt.score(X_test_fs, y_test))
```

0. 823708206687

In [35]:

```
#####  
#  
#           Regularization  
#  
#####  
  
#linearregression  
#####  
#>>> 1> Input dataset  
X_train=[[6], [8], [10], [14], [18]]  
y_train=[[7], [9], [13], [17.5], [18]]  
  
#####  
#>>> 2> Build linearregression model  
from sklearn.linear_model import LinearRegression  
regressor=LinearRegression()  
regressor.fit(X_train,y_train)  
  
#####  
#>>> 3> Predict  
xx=np.linspace(0, 26, 100)  
xx=xx.reshape(xx.shape[0], 1)  
yy=regressor.predict(xx)  
  
#####  
#>>> 4> Visualization and Output  
import matplotlib.pyplot as plt  
  
plt.scatter(X_train,y_train)  
plt1,=plt.plot(xx, yy, label=' Degree=1')  
plt.axis([0, 25, 0, 25])  
plt.xlabel(' Diameter of Pizza')  
plt.ylabel(' Price of Pizza')  
plt.legend(handles=[plt1])  
plt.show()  
  
print('The R-squared value of Linear Regressor performing on the training data is',regressor.score(X
```



The R-squared value of Linear Regressor performing on the training data is 0.910001596424

In [36]:

```
#polynomial featrues in degree 2
#####
#>>> 2&3> Build model and Predict
from sklearn.preprocessing import PolynomialFeatures


poly2=PolynomialFeatures(degree=2)
X_train_poly2=poly2.fit_transform(X_train)

regressor_poly2=LinearRegression()
regressor_poly2.fit(X_train_poly2,y_train)

xx_poly2=poly2.transform(xx)
yy_poly2=regressor_poly2.predict(xx_poly2)

#####
#>>> 4> Visualization and Output
plt.scatter(X_train,y_train)
plt1,=plt.plot(xx,yy,label='Degree=1')
plt2,=plt.plot(xx,yy_poly2,label='Degree=2')
plt.axis([0,25,0,25])
plt.xlabel('Diameter of Pizza')
plt.ylabel('Price of Pizza')
plt.legend(handles=[plt1,plt2])
plt.show()

print('The R-squared value of Polynomial Regressor(Degree=2) performing on the training data is',re
```



The R-squared value of Polynomial Regressor(Degree=2) performing on the training data is 0.98164216396

In [39]:

```
#polynomial featrues in degree 4
#####
#>>> 2&3> Build model and Predict
from sklearn.preprocessing import PolynomialFeatures

poly4=PolynomialFeatures(degree=4)
X_train_poly4=poly4.fit_transform(X_train)

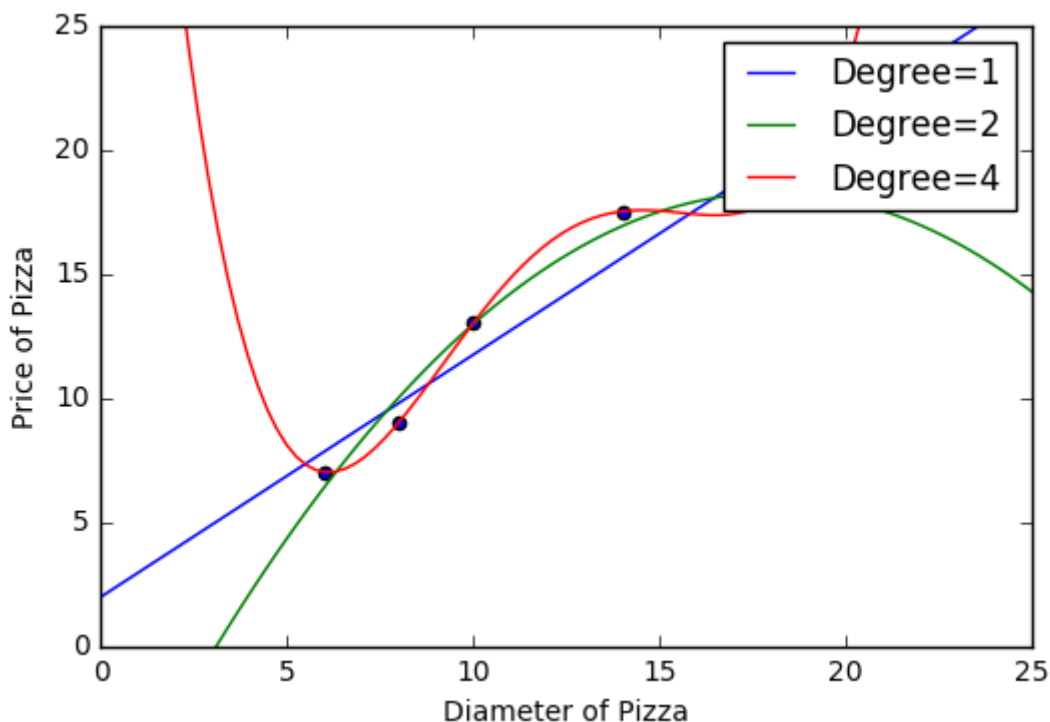
regressor_poly4=LinearRegression()
regressor_poly4.fit(X_train_poly4,y_train)

xx_poly4=poly4.transform(xx)
yy_poly4=regressor_poly4.predict(xx_poly4)
print(regressor_poly4.coef_)
print(np.sum(regressor_poly4.coef_**2))

#####
#>>> 4> Visualization and Output
plt.scatter(X_train,y_train)
plt1,=plt.plot(xx,yy,label='Degree=1')
plt2,=plt.plot(xx,yy_poly2,label='Degree=2')
plt4,=plt.plot(xx,yy_poly4,label='Degree=4')
plt.axis([0,25,0,25])
plt.xlabel('Diameter of Pizza')
plt.ylabel('Price of Pizza')
plt.legend(handles=[plt1,plt2,plt4])
plt.show()

print('The R-squared value of Polynominal Regressor(Degree=4) performing on the training data is',re
```

```
[[ 0.00000000e+00 -2.51739583e+01  3.68906250e+00 -2.12760417e-01
  4.29687500e-03]]
647.382645737
```



The R-squared value of Polynominal Regressor(Degree=4) performing on the training data is 1.0

In [40]:

```
#####  
#>>> 5> Validation  
  
#Prepare test dataset  
X_test=[[6],[8],[11],[16]]  
y_test=[[8],[12],[15],[18]]  
  
#linear regression  
print('The R-squared value of Linear Regressor performing on the test data is',regressor.score(X_test,y_test))  
  
#polynomial 2  
X_test_poly2=poly2.fit_transform(X_test)  
print('The R-squared value of Polynomial Regressor(Degree=2) performing on the test data is',regressor.score(X_test_poly2,y_test))  
  
#polynomial 4  
X_test_poly4=poly4.fit_transform(X_test)  
print('The R-squared value of Polynomial Regressor(Degree=4) performing on the test data is',regressor.score(X_test_poly4,y_test))
```

The R-squared value of Linear Regressor performing on the test data is 0.80972683246

7

The R-squared value of Polynomial Regressor(Degree=2) performing on the test data is 0.867544365635

The R-squared value of Polynomial Regressor(Degree=4) performing on the test data is 0.809588079577

In [42]:

```
#lasso  
#####  
#>>> 2&3> Build model and Predict  
from sklearn.linear_model import Lasso  
  
lasso_poly4=Lasso()  
lasso_poly4.fit(X_train_poly4,y_train)  
print(lasso_poly4.score(X_test_poly4,y_test))  
print(lasso_poly4.coef_)
```

0.83889268736

```
[ 0.00000000e+00  0.00000000e+00  1.17900534e-01  5.42646770e-05  
 -2.23027128e-04]
```

C:\Program Files\Anaconda3\lib\site-packages\sklearn\linear_model\coordinate_descent.py:466: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations
ConvergenceWarning)